

## Structure And Design

---

First, I designed a structure called *buffer\_t*, which contains the data to be sent/ acknowledged/ submitted. Also, some helper verities like *Length*, *usedLength*, *lastSentTime* are also in this structure.

Each connection has three *buffer\_t* linked list to manage the data:

- *unsentList*: Contains the data from the input, but hasn't been transferred to the peer.
- *sentUnackList*: Contains the data which has been sent but not been acknowledged. Maybe needed to resend in the future.
- *unsubmittedList*: Contains the data gotten from the sender, but hasn't been submitted to the output.

Here we should notice: Considering without selective ack option, **we will not need to visit the *buffer\_t* randomly** (except on-timer check). We only pop out the buffers from the front and insert new buffers at the tail, where we greatly reduce the time complexity.

We also should know that **not every movement will lead to a buffer moved to another list**. For example, a buffer with 500Byte data can only send 200Byte at a time. So, we also need to record *usedLength* of the buffer, and *memcpy()* the sent data to *sentUnackList*, but buffers in *sentUnackList* and *unsubmittedList* are handled as a whole structure.

## Process

---

### Sending

---

*Send()* should not only be triggered by *input()*. For example, when it receives the data that cannot be used or the other side announces the zero window, or there has been a lot of data in the *unsentList*. So, the better way is also triggering the *send()* on timer and when receiving anything. But with this circumstance, *send()* may have nothing to send when triggered. So, I named it as *trySend()* -- if it has nothing to send, just return.

*trySend* has three job: sending data, sending ack and sending FIN, and they can be compatible at a time.

- Sending data: calculate how much it can send from total data in the *unsentList*, peer's *rwnd*, it's *cwnd*, and choose the minimum value and prepare data.
- Ack number: get the latest ack number from global state's variable.
- FIN: decided by connection's status and the left data in *unsentList*. If it cannot send all left data this time, it cannot send FIN.

The last thing to do is hanging the buffer to the *sentUnackList*.

### Receiving

---

When it gets the data, it should validate the packet, like the *checksum*, the real length and if the *seqNo* equal than current record. Then, it determines how much data can be acked from the *ackNo*, and pop out that much buffer(s) from the *sentUnackList* and hang it/them to *unsubmittedList*.

When it gets some "useless" packets, like outdated segments or 0 data size segment or unexpected seqNo, it means the sender needs the help from the receiver to synchronize, so it should immediately send a ack to the sender.

## Timer

---

When on timer, it checks all buffers in *sentUnackList*, and decide whether to resend it, and whether to shut down the connection.

## Challenges

---

1. Handling EOF: In the initial version, when the receiver saw the FIN, it immediately close the output channel. After analyze byte by byte of the output file and trace log. I found only the data when piggybacking with the FIN will be lost. Then, I realize I should lower the priority of handling FIN on receiver side.
2. Zero window problem: When testing with a slow outputting receiver, it could be possible the receiver announces a zero rwnd to the sender, and the sender will never send anything. And because the receiver gets nothing, it also won't update the rwnd to the sender. So, I tried to modify the code -- the sender will periodically send meaningless segment -- a segment with no data. But here is another question, the receiver won't reply anything if nothing to update to the sender. So, I forced the receiver to reply once it gets a segment.
3. Avoid too many meaningless segments in flight. Because the sender and receiver will both send empty segments, these segments may waste too much resource. So, I set a flag: only when it has ack to announce, or has data to send, or meets zero-window problem, it can send dataless segments.

## Tests

---

- Test on `ctcp_tests.py` : score from 13 to 17.  
I tried to find out if a larger/ smaller timeout can get a better result, but it's clueless.
- Test on `generic_tester.py` `ctcp_mininet.xml` : score 240/240
- Test on `generic_tester.py` `ctcp.xml` : score 40/40  
Spent most of time on it. It turns out that I need to change the timeout punishment from 2 RTT to 5 RTT.
- Test on Internet Web servers : sucessful.

## Bugs

---

Currently, it uses the config's RTT time and never update it. I think a good TCP model should measure the RTT all the time. Also, there should be a notice when the sender reads the EOF, otherwise we cannot know if it has sent over all the data and keep waiting.